

Code :

```
package SQL

import org.apache.spark.sql.SparkSession

object SparkSQLUseCase1 {

  case class
  hvac_cls(Date:String,Time:String,TargetTemp:Int,ActualTemp:Int,System:Int,SystemAge
:Int,BuildingId:Int)
  //For implicit conversions like converting RDDs and sequences to DataFrames
  case class
  building(buildid:Int,buildmgr:String,buildAge:Int,hvacproduct:String,Country:String
)

  def main(args: Array[String]): Unit = {
    println("hey scala") //Let us create a spark session object

    val spark = SparkSession
      .builder()
      .master("local")
      .appName("Spark SQL Use Case 1 ")
      .config("spark.some.config.option", "some-value")
      .getOrCreate()
    println("Spark Session Object created")

    val data =
    spark.sparkContext.textFile("C:/Users/admin/Desktop/bigdata/Spark/Spark SQL Case
Studies/HVAC.csv");
    println("HVAC Data->>" + data.count())

    val header = data.first()
    val data1 = data.filter(row => row != header)
    println("Header removed from the data !")

    //For implicit conversions like converting RDDs and sequences to DataFrames
    import spark.implicits._
    val hvac = data1.map(x=>x.split(",")).map(x =>
    hvac_cls(x(0),x(1),x(2).toInt,x(3).toInt,x(4).toInt,x(5).toInt,x(6).toInt)).toDF()
    hvac.show()
    println("HVAC Dataframe created !")

    hvac.registerTempTable("HVAC")
    println("Dataframe Registered as table !")
    //createOrReplaceTempView

    val hvac1 = spark.sql("select *,IF((targettemp - actualtemp) > 5, '1',
    IF((targettemp - actualtemp) < -5, '1', 0)) AS tempchange from HVAC")
    hvac1.show()
    hvac1.registerTempTable("HVAC1")
    println("Data Frame Registered as HVAC1 table !")

    //Now lets load the second data set
    val data2 =
    spark.sparkContext.textFile("C:/Users/admin/Desktop/bigdata/Spark/Spark SQL Case
Studies/building.csv");
    val header1 = data2.first()
    val data3 = data2.filter(row => row != header1)
    println("Header removed from the building data")
    println("Buildings Data->>" + data3.count())

    //Now let us create the building dataframe
    val build = data3.map(x=> x.split(",")).map(x =>
```

```

building(x(0).toInt,x(1),x(2).toInt,x(3),x(4)).toDF
  build.show()
  build.registerTempTable("building")
  println("Buildings data registered as building table")

  //Now join the two tables
  val build1 = spark.sql("select h.*, b.country, b.hvacproduct from building b
join hvac1 h on b.buildid = h.buildingid")

  build1.show()

  //Select temperature and country column from above

  val tempCountry = build1.map(x => (new Integer(x(7).toString),x(8).toString))

  tempCountry.show()

  //Filter the values

  val tempCountryOnes = tempCountry.filter(x=> {if(x._1==1) true else false})

  tempCountryOnes.show()

  tempCountryOnes.groupBy("_2").count.show

  //Save the output to the disk
  //
tempCountryOnes.write.save("C:/Users/admin/Desktop/bigdata/Spark/sparksqloutput")

}

}

```

Output:

```

18/09/12 21:42:59 INFO TaskSchedulerImpl: Removed TaskSet 0.0, whose tasks have all completed, from pool
18/09/12 21:42:59 INFO DAGScheduler: ResultStage 0 (count at SparkSQLUseCase1.scala:24) finished in 0.417 s
HVAC Data->>8001
18/09/12 21:43:00 INFO DAGScheduler: Job 0 finished: count at SparkSQLUseCase1.scala:24, took 0.661743 s
18/09/12 21:43:00 INFO SparkContext: Starting job: first at SparkSQLUseCase1.scala:26

18/09/12 21:43:00 INFO DAGScheduler: ResultStage 1 (first at SparkSQLUseCase1.scala:26) finished in 0.036 s
18/09/12 21:43:00 INFO DAGScheduler: Job 1 finished: first at SparkSQLUseCase1.scala:26, took 0.047574 s
Header removed from the data !
18/09/12 21:43:00 INFO BlockManagerInfo: Removed broadcast_2_piece0 on 192.168.100.3:62159 in memory (size: 2002.0 B, free: 897.6 MB)
18/09/12 21:43:05 INFO CodeGenerator: Code generated in 429.637203 ms

```

```
18/09/12 21:43:05 INFO DAGScheduler: Job 2 finished: show at SparkSQLUseCase1.scala:33, took 0.063892 s
18/09/12 21:43:05 INFO CodeGenerator: Code generated in 38.333182 ms
```

Date	Time	TargetTemp	ActualTemp	System	SystemAge	BuildingId
6/1/13	0:00:01	66	58	13	20	4
6/2/13	1:00:01	69	68	3	20	17
6/3/13	2:00:01	70	73	17	20	18
6/4/13	3:00:01	67	63	2	23	15
6/5/13	4:00:01	68	74	16	9	3
6/6/13	5:00:01	67	56	13	28	4
6/7/13	6:00:01	70	58	12	24	2
6/8/13	7:00:01	70	73	20	26	16
6/9/13	8:00:01	66	69	16	9	9
6/10/13	9:00:01	65	57	6	5	12
6/11/13	10:00:01	67	70	10	17	15
6/12/13	11:00:01	69	62	2	11	7
6/13/13	12:00:01	69	73	14	2	15
6/14/13	13:00:01	65	61	3	2	6
6/15/13	14:00:01	67	59	19	22	20
6/16/13	15:00:01	65	56	19	11	8
6/17/13	16:00:01	67	57	15	7	6
6/18/13	17:00:01	66	57	12	5	13
6/19/13	18:00:01	69	58	8	22	4
6/20/13	19:00:01	67	55	17	5	7

only showing top 20 rows

HVAC Dataframe created !

```
18/09/12 21:43:05 INFO BlockManagerInfo: Removed broadcast_1_piece0 on 192.168.100.3:62159 in memory (size: 1925.0 B, free: 897.6 MB)
```

HVAC Dataframe created !

```
18/09/12 21:43:05 INFO BlockManagerInfo: Removed broadcast_1_piece0 on 192.168.100.3:62159 in memory (size: 1925.0 B, free: 897.6 MB)
```

```
18/09/12 21:43:05 INFO BlockManagerInfo: Removed broadcast_3_piece0 on 192.168.100.3:62159 in memory (size: 5.2 KB, free: 897.6 MB)
```

```
18/09/12 21:43:05 INFO ContextCleaner: Cleaned accumulator 96
```

```
18/09/12 21:43:05 INFO ContextCleaner: Cleaned accumulator 97
```

```
18/09/12 21:43:05 INFO SparkSqlParser: Parsing command: HVAC
```

Dataframe Registered as table !

```
18/09/12 21:43:06 INFO SparkSqlParser: Parsing command: select *,IF((targettemp - actualtemp) > 5, '1', IF((targettemp - actualtemp) < -5, '1', 0)) AS tempchange from HVAC
```

```
18/09/12 21:43:06 INFO CodeGenerator: Code generated in 142.64447 ms
```

```
18/09/12 21:43:06 INFO SparkContext: Starting job: show at SparkSQLUseCase1.scala:41
```

```
18/09/12 21:43:06 INFO CodeGenerator: Code generated in 28.968618 ms
```

```
18/09/12 21:43:06 INFO SparkSqlParser: Parsing command: HVAC1
```

Date	Time	TargetTemp	ActualTemp	System	SystemAge	BuildingId	tempchange
6/1/13	0:00:01	66	58	13	20	4	1
6/2/13	1:00:01	69	68	3	20	17	0
6/3/13	2:00:01	70	73	17	20	18	0
6/4/13	3:00:01	67	63	2	23	15	0
6/5/13	4:00:01	68	74	16	9	3	1
6/6/13	5:00:01	67	56	13	28	4	1
6/7/13	6:00:01	70	58	12	24	2	1
6/8/13	7:00:01	70	73	20	26	16	0
6/9/13	8:00:01	66	69	16	9	9	0
6/10/13	9:00:01	65	57	6	5	12	1
6/11/13	10:00:01	67	70	10	17	15	0
6/12/13	11:00:01	69	62	2	11	7	1
6/13/13	12:00:01	69	73	14	2	15	0
6/14/13	13:00:01	65	61	3	2	6	0
6/15/13	14:00:01	67	59	19	22	20	1
6/16/13	15:00:01	65	56	19	11	8	1
6/17/13	16:00:01	67	57	15	7	6	1
6/18/13	17:00:01	66	57	12	5	13	1
6/19/13	18:00:01	69	58	8	22	4	1
6/20/13	19:00:01	67	55	17	5	7	1

only showing top 20 rows

Data Frame Registered as HVAC1 table !

```
18/09/12 21:43:06 INFO MemoryStore: Block broadcast_5 stored as values in memory (estimated size 127.2 KB, free 897.3 MB)
```

```
18/09/12 21:43:06 INFO DAGScheduler: Job 4 finished: first at SparkSQLUseCase1.scala:47, took 0.046644 s
Header removed from the building data
18/09/12 21:43:06 INFO SparkContext: Starting job: count at SparkSQLUseCase1.scala:50
18/09/12 21:43:06 INFO DAGScheduler: Got job 5 (count at SparkSQLUseCase1.scala:50) with 1 output partitions
18/09/12 21:43:06 INFO DAGScheduler: Final stage: ResultStage 5 (count at SparkSQLUseCase1.scala:50)
18/09/12 21:43:06 INFO DAGScheduler: Parents of final stage: List()
18/09/12 21:43:06 INFO DAGScheduler: Missing parents: List()
18/09/12 21:43:06 INFO DAGScheduler: Submitting ResultStage 5 (MapPartitionsRDD[15] at filter at SparkSQLUseCase1.scala:48), which has no mi
18/09/12 21:43:06 INFO MemoryStore: Block broadcast_7 stored as values in memory (estimated size 3.4 KB, free 897.3 MB)
18/09/12 21:43:06 INFO MemoryStore: Block broadcast_7_piece0 stored as bytes in memory (estimated size 2.0 KB, free 897.3 MB)
18/09/12 21:43:06 INFO BlockManagerInfo: Added broadcast_7_piece0 in memory on 192.168.100.3:62159 (size: 2.0 KB, free: 897.6 MB)
18/09/12 21:43:06 INFO SparkContext: Created broadcast 7 from broadcast at DAGScheduler.scala:996
18/09/12 21:43:06 INFO DAGScheduler: Submitting 1 missing tasks from ResultStage 5 (MapPartitionsRDD[15] at filter at SparkSQLUseCase1.scala:
18/09/12 21:43:06 INFO TaskSchedulerImpl: Adding task set 5.0 with 1 tasks
18/09/12 21:43:06 INFO TaskSetManager: Starting task 0.0 in stage 5.0 (TID 5, localhost, executor driver, partition 0, PROCESS_LOCAL, 5942 b
18/09/12 21:43:06 INFO Executor: Running task 0.0 in stage 5.0 (TID 5)
18/09/12 21:43:06 INFO HadoopRDD: Input split: file:/C:/Users/admin/Desktop/bigdata/Spark/Spark SQL Case Studies/building.csv:0+544
Buildings Data-->20
18/09/12 21:43:06 INFO Executor: Finished task 0.0 in stage 5.0 (TID 5). 946 bytes result sent to driver
```

```
18/09/12 21:43:07 INFO SparkSqlParser: Parsing command: building
+-----+
|buildid|buildmgr|buildAge|hvacproduct|Country|
+-----+
| 1| M1| 25| AC1000| USA|
| 2| M2| 27| FN39TG| France|
| 3| M3| 28| JDN577| Brazil|
| 4| M4| 17| GGI919| Finland|
| 5| M5| 3| ACMA22| Hong Kong|
| 6| M6| 9| AC1000| Singapore|
| 7| M7| 13| FN39TG| South Africa|
| 8| M8| 25| JDN577| Australia|
| 9| M9| 11| GGI919| Mexico|
| 10| M10| 23| ACMA22| China|
| 11| M11| 14| AC1000| Belgium|
| 12| M12| 26| FN39TG| Finland|
| 13| M13| 25| JDN577| Saudi Arabia|
| 14| M14| 17| GGI919| Germany|
| 15| M15| 19| ACMA22| Israel|
| 16| M16| 23| AC1000| Turkey|
| 17| M17| 11| FN39TG| Egypt|
| 18| M18| 25| JDN577| Indonesia|
| 19| M19| 14| GGI919| Canada|
| 20| M20| 19| ACMA22| Argentina|
+-----+

Buildings data registered as building table
18/09/12 21:43:07 INFO SparkSqlParser: Parsing command: select h.*, b.country, b.hvacproduct from building b join hvac1 h on b.buildid = h.buildingid
18/09/12 21:43:07 INFO ContextCleaner: Cleaned accumulator 146
```

```
18/09/12 21:43:09 INFO ContextCleaner: Cleaned accumulator 342
18/09/12 21:43:09 INFO ContextCleaner: Cleaned accumulator 343
+-----+
|Date|Time|TargetTemp|ActualTemp|System|SystemAge|BuildingId|tempchange|country|hvacproduct|
+-----+
|6/10/13|9:00:01|65|57|6|5|12|1|Finland|FN39TG|
|6/18/13|23:13:19|66|75|1|13|12|1|Finland|FN39TG|
|6/2/13|13:43:51|65|72|20|26|12|1|Finland|FN39TG|
|6/13/13|0:13:20|67|77|8|19|12|1|Finland|FN39TG|
|6/16/13|3:13:20|67|55|11|16|12|1|Finland|FN39TG|
|6/30/13|17:13:20|65|57|17|9|12|1|Finland|FN39TG|
|6/1/13|18:13:20|68|65|7|21|12|0|Finland|FN39TG|
|6/25/13|18:33:07|70|66|20|20|12|0|Finland|FN39TG|
|6/17/13|16:00:01|69|68|16|4|12|0|Finland|FN39TG|
|6/5/13|16:43:51|69|69|19|15|12|0|Finland|FN39TG|
|6/23/13|10:13:20|65|61|1|1|12|0|Finland|FN39TG|
|6/29/13|16:13:20|67|80|12|8|12|1|Finland|FN39TG|
|6/4/13|21:13:20|66|72|7|1|12|1|Finland|FN39TG|
|6/3/13|2:00:01|69|72|7|21|12|0|Finland|FN39TG|
|6/16/13|15:00:01|67|77|4|22|12|1|Finland|FN39TG|
|6/22/13|21:00:01|70|77|13|12|12|1|Finland|FN39TG|
|6/26/13|7:43:51|65|62|6|6|12|0|Finland|FN39TG|
|6/26/13|13:13:20|65|63|20|9|12|0|Finland|FN39TG|
|6/30/13|17:13:20|66|62|14|26|12|0|Finland|FN39TG|
|6/10/13|3:33:07|70|78|5|9|12|1|Finland|FN39TG|
+-----+
only showing top 20 rows
```

```
18/09/12 21:43:09 INFO BlockManagerInfo: Removed broadcast_9_piece0 on 192.168.100.3:62159 in memory (size: 7.4 KB, free: 897.6 MB)
18/09/12 21:43:10 INFO CodeGenerator: Code generated in 102.665842 ms
```

18/09/12 21:43:11 INFO CodeGenerator: Code generated in 8.403279 ms

```
+-----+
| _1|   _2|
+-----+
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 0|Finland|
| 0|Finland|
| 0|Finland|
| 0|Finland|
| 0|Finland|
| 1|Finland|
| 1|Finland|
| 0|Finland|
| 1|Finland|
| 1|Finland|
| 0|Finland|
| 0|Finland|
| 0|Finland|
| 1|Finland|
+-----+
```

only showing top 20 rows

18/09/12 21:43:11 INFO ContextCleaner: Cleaned accumulator 1130

18/09/12 21:43:11 INFO ContextCleaner: Cleaned accumulator 1131

18/09/12 21:43:12 INFO ShuffleBlockFetcherIterator: Started 0 remote fetches in 0 ms

```
+-----+
| _1|   _2|
+-----+
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
| 1|Finland|
+-----+
```

only showing top 20 rows

18/09/12 21:43:12 INFO Executor: Finished task 19.0 in stage 33.0 (TID 87). 3163 bytes result sent to driver

18/09/12 21:43:18 INFO DAGScheduler: ResultStage 53 (show at SparkSQLUseCase1.scala:75) finished in 0.347 s

18/09/12 21:43:18 INFO DAGScheduler: Job 20 finished: show at SparkSQLUseCase1.scala:75, took 0.354965 s

```
+-----+
|      _2|count|
+-----+
| Singapore| 230|
| Turkey| 243|
| Germany| 196|
| France| 251|
| Argentina| 230|
| Belgium| 199|
| Finland| 473|
| China| 241|
| Hong Kong| 248|
| Israel| 232|
| USA| 213|
| Mexico| 228|
| Indonesia| 243|
| Saudi Arabia| 233|
| Canada| 232|
| Brazil| 226|
| Australia| 225|
| Egypt| 236|
| South Africa| 237|
+-----+
```

18/09/12 21:43:18 INFO CodeGenerator: Code generated in 9.042765 ms

18/09/12 21:43:18 INFO SparkContext: Invoking stop() from shutdown hook
